

Chapter 15

RNA-Seq Data

15.1 Introduction

The limma approach to RNA-seq explained in the articles by Law et al [15] and Liu et al [20]. See also the article by Law et al [14], which gives a complete workflow case study. In the limma approach to RNA-seq, read counts are converted to log2-counts-per-million (logCPM) and the mean-variance relationship is modelled either with precision weights or with an empirical Bayes prior trend. The precision weights approach is called “voom” and the prior trend approach is called “limma-trend” [15]. In either case, the RNA-seq data can be analyzed as if it was microarray data. This means that any of the linear modelling or gene set testing methods in the limma package can be applied to RNA-seq data.

15.2 Making a count matrix

RNA-seq data usually arrives in the form of FastQ or BAM files of unaligned reads. The reads need to be mapped to a reference genome or transcriptome, then summarized at the exon or gene level to produce a matrix of counts. We find the Rsubread package [17] to be convenient, fast and effective for this purpose. Other popular methods include RSEM [16] and HTseq. A runnable example with complete code showing how to use Rsubread with limma is provided at <http://bioinf.wehi.edu.au/RNAseqCaseStudy>. Another complete code example is provided by Chen et al [6].

15.3 Normalization and filtering

Once a matrix of read counts `counts` has been created, with rows for genes and columns for samples, it is convenient to create a `DGEList` object using the edgeR package:

```
> dge <- DGEList(counts=counts)
```

The next step is to remove rows that consistently have zero or very low counts. One can for example use

```
> keep <- filterByExpr(dge, design)
> dge <- dge[keep,,keep.lib.sizes=FALSE]
```

where `filterByExpr` is a function in the edgeR package. Here we will assume that filtering has been done.

It is usual to apply scale normalization to RNA-seq read counts, and the TMM normalization method [33] in particular has been found to perform well in comparative studies. This can be applied to the `DGEList` object:

```
> dge <- calcNormFactors(dge)
```

15.4 Differential expression: limma-trend

If the sequencing depth is reasonably consistent across the RNA samples, then the simplest and most robust approach to differential exis to use `limma-trend`. This approach will usually work well if the ratio of the largest library size to the smallest is not more than about 3-fold.

In the `limma-trend` approach, the counts are converted to `logCPM` values using `edgeR`'s `cpm` function:

```
> logCPM <- cpm(dge, log=TRUE, prior.count=3)
```

The prior count is used here to damp down the variances of logarithms of low counts.

The `logCPM` values can then be used in any standard `limma` pipeline, using the `trend=TRUE` argument when running `eBayes` or `treat`. For example:

```
> fit <- lmFit(logCPM, design)
> fit <- eBayes(fit, trend=TRUE)
> topTable(fit, coef=ncol(design))
```

Or, to give more weight to fold-changes in the gene ranking, one might use:

```
> fit <- lmFit(logCPM, design)
> fit <- treat(fit, lfc=log2(1.2), trend=TRUE)
> topTreat(fit, coef=ncol(design))
```

15.5 Differential expression: voom

When the library sizes are quite variable between samples, then the `voom` approach is theoretically more powerful than `limma-trend`. In this approach, the `voom` transformation is applied to the normalized and filtered `DGEList` object:

```
v <- voom(dge, design, plot=TRUE)
```

The `voom` transformation uses the experiment design matrix, and produces an `EList` object.

It is also possible to give a matrix of counts directly to `voom` without TMM normalization, by

```
> v <- voom(counts, design, plot=TRUE)
```

If the data are very noisy, one can apply the same between-array normalization methods as would be used for microarrays, for example:

```
> v <- voom(counts, design, plot=TRUE, normalize="quantile")
```

After this, the usual `limma` pipelines for differential expression can be applied, for example:

```
> fit <- lmFit(v, design)
> fit <- eBayes(fit)
> topTable(fit, coef=ncol(design))
```

Or, to give more weight to fold-changes in the ranking, one could use say:


```
> fit <- treat(fit, lfc=log2(1.2))
> topTreat(fit, coef=ncol(design))
```

15.6 Voom with sample quality weights

When a multi-dimensional scaling plot from a designed RNA-seq experiment indicates the presence of outlier samples, it is possible to combine the observation-level weighting strategy used in voom with sample-specific quality weights (as described in the section above on Array Quality Weights) to down-weight outlier samples. This capability is implemented in the `voomWithQualityWeights` function.

The example below shows its use on an RNA-seq data set where the epigenetic regulator *Smchd1* has been knocked-out in lymphoma cell-lines (GEO series GSE64099) [20]. Overall we obtain more differential expression by applying this combined weighting strategy and the raw *p*-value and false discovery rate for the *Smchd1* gene, which has been knocked out, is smaller.

```
> plotMDS(x, labels=1:7, col=as.numeric(genotype), main="MDS plot")
> legend("topright", legend=c("WT", "KO"), col=1:2, pch=15)
```




```
> # Analysis with voom only
> des[1:7,]
      (Intercept) Smchd1nullvsWt
1             1             1
2             1             1
3             1             1
4             1             1
5             1             0
6             1             0
7             1             0
> v <- voom(x, design=des)
> plotMDS(v, labels=1:7, col=as.numeric(genotype))
> vfit <- lmFit(v)
> vfit <- eBayes(vfit)
> options(digits=3)
> topTable(vfit,coef=2,sort.by="P")
      GeneID      Symbols logFC AveExpr      t P.Value adj.P.Val      B
74355  74355      Smchd1 -3.12   6.067 -23.35 2.16e-08 0.000266 9.97
```

```

18028 18028 Nfib 8.98 1.714 12.60 2.17e-06 0.013355 3.15
75605 75605 Kdm5b -3.55 3.618 -11.75 3.62e-06 0.014857 5.06
667435 667435 Igkv17-121 -5.35 -1.435 -10.22 9.95e-06 0.025513 2.57
381126 381126 Gareme 6.17 0.113 10.08 1.10e-05 0.025513 2.35
381413 381413 Gpr176 -4.02 1.328 -9.90 1.25e-05 0.025513 3.39
75033 75033 Mei4 6.44 0.259 9.69 1.45e-05 0.025513 2.23
69136 69136 Tusc1 5.67 -0.184 8.90 2.67e-05 0.040995 1.87
233552 233552 Gdpd5 -2.82 1.948 -8.56 3.49e-05 0.042754 2.81
80890 80890 Trim2 -1.43 4.491 -8.40 4.00e-05 0.042754 2.72
> top <- topTable(vfit,coef=2,number=Inf,sort.by="P")
> sum(top$adj.P.Val<0.05)
[1] 12
> # Analysis with combined voom and sample quality weights
> vwts <- voomWithQualityWeights(x, design=des, normalization="none", plot=TRUE)
> vfit2 <- lmFit(vwts)
> vfit2 <- eBayes(vfit2)
> top2 <- topTable(vfit2,coef=2,sort.by="P")
      GeneID Symbols logFC AveExpr      t P.Value adj.P.Val      B
74355 74355 Smchd1 -3.17 6.067 -28.5 1.61e-09 1.98e-05 12.57
18028 18028 Nfib 9.23 1.714 19.0 4.44e-08 2.73e-04 6.91
381126 381126 Gareme 6.45 0.113 15.9 1.85e-07 7.58e-04 6.02
75033 75033 Mei4 6.56 0.259 15.0 2.84e-07 8.73e-04 5.83
69136 69136 Tusc1 5.88 -0.184 13.6 6.16e-07 1.11e-03 5.31
54354 54354 Rassf5 5.74 4.554 13.6 6.26e-07 1.11e-03 6.63
75605 75605 Kdm5b -3.80 3.618 -13.5 6.53e-07 1.11e-03 6.67
58998 58998 Pvr13 7.69 0.961 13.1 8.46e-07 1.11e-03 5.33
320398 320398 Lrig3 7.39 1.584 13.1 8.49e-07 1.11e-03 5.32
17069 17069 Ly6e 2.63 7.605 13.0 9.01e-07 1.11e-03 6.26
> top2 <- topTable(vfit2,coef=2,number=Inf,sort.by="P")
> sum(top2$adj.P.Val<0.05)
[1] 1478

```



15.7 Differential splicing

limma can also detect genes that show evidence of differential splicing between conditions. One can test for differential splicing associated with any contrast for a linear model.

In this case, the matrix of counts should be at the exon level, with a row for each exon. For example,

```
> dge <- DGEList(counts=counts)
> dge$genes$GeneID <- GeneID
```

where `counts` is a matrix of exon-level counts, and `GeneID` identifies which gene each exon belongs to. Then filter and normalize:

```
> A <- rowSums(dge$counts)
> dge <- dge[A>10,, keep.lib.sizes=FALSE]
> dge <- calcNormFactors(dge)
```

Then apply the voom transformation and fit a linear model:

```
> v <- voom(dge, design, plot=TRUE)
> fit <- lmFit(v, design)
```

Now we can test for differential splicing associated with any coefficient in the linear model. First run the `diffSplice` function:

```
> ex <- diffSplice(fit, geneid="GeneID")
```

Then

```
> topSplice(ex, coef=2, test="simes")
```

will find genes that show evidence of differential splicing associated with the second coefficient in the linear model. The output is similar that from the limma `topTable` function. More detail can be obtained by

```
> topSplice(ex, coef=2, test="t")
```

which will show individual exons that are enriched or depleted relative to other exons in the same gene. To display the pattern of exons in the top genes:

```
> plotSplice(ex)
```